

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.colors as mcolors
from scipy import stats
import seaborn as sns
from textwrap import wrap
from sklearn.cluster import KMeans
from sklearn.manifold import TSNE
from sklearn.preprocessing import StandardScaler

import warnings
warnings.filterwarnings('ignore')

#Reading train and test dataset in pandas dataframe
train_df = pd.read_excel('train.xlsx')
test_df = pd.read_excel('test.xlsx')

# Concatenate train & test dataset
df = pd.concat([train_df,test_df], axis=0)

print("Dimension of the Data set is", df.shape)
print()
df.head()

```

Dimension of the Data set is (541909, 8)

	InvoiceNo	StockCode	Description	Quantity	\
0	558904	22292	HANGING CHICK YELLOW DECORATION	1	
1	556072	20970	PINK FLORAL FELTCRAFT SHOULDER BAG	8	
2	551739	21559	STRAWBERRY LUNCH BOX WITH CUTLERY	2	
3	541658	21988	PACK OF 6 SKULL PAPER PLATES	1	
4	538364	85099C	JUMBO BAG BAROQUE BLACK WHITE	10	

	InvoiceDate	UnitPrice	CustomerID	Country
0	2011-07-04 16:18:00	1.25	NaN	United Kingdom
1	2011-06-08 14:57:00	3.75	16126.0	United Kingdom
2	2011-05-04 10:58:00	2.55	18118.0	United Kingdom
3	2011-01-20 12:16:00	0.85	15529.0	United Kingdom
4	2010-12-10 17:26:00	1.95	14448.0	United Kingdom

df.info()

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 541909 entries, 0 to 162572
Data columns (total 8 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   InvoiceNo        541909 non-null   object 
 1   StockCode         541909 non-null   object 

```

```

2   Description  540455 non-null  object
3   Quantity     541909 non-null  int64
4   InvoiceDate  541909 non-null  datetime64[ns]
5   UnitPrice    541909 non-null  float64
6   CustomerID   406829 non-null  float64
7   Country      541909 non-null  object
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 37.2+ MB

```

```
df.describe().transpose()
```

	count	mean	std	min	25%
50% \					
Quantity	541909.0	9.552250	218.081158	-80995.00	1.00
3.00					
UnitPrice	541909.0	4.611114	96.759853	-11062.06	1.25
2.08					
CustomerID	406829.0	15287.690570	1713.600303	12346.00	13953.00
15152.00					
		75%	max		
Quantity		10.00	80995.0		
UnitPrice		4.13	38970.0		
CustomerID	16791.00	18287.0			

```
# Checking all the records which were returned
df[df["InvoiceNo"].str.startswith('C',na=False)]
```

	InvoiceNo	StockCode	Description
Quantity \			
415	C549269	75049L	LARGE CIRCULAR MIRROR MOBILE
-9			
487	C572109	23064	CINDERELLA CHANDELIER
-1			
613	C537860	22180	RETROSPOT LAMP
-1			
834	C560540	23240	SET OF 4 KNICK KNACK TINS DOILEY
-1			
874	C542910	20726	LUNCH BAG WOODLAND
-1			
...	...	...	...
162302	C558712	21735	TWO DOOR CURIO CABINET
-1			
162334	C550780	84507C	BLUE CIRCLES DESIGN MONKEY DOLL
-1			
162344	C553031	21533	RETROSPOT LARGE MILK JUG
-3			
162421	C542910	85123A	WHITE HANGING HEART T-LIGHT HOLDER
-1			
162458	C543368	22941	CHRISTMAS LIGHTS 10 REINDEER

- 4

```
    InvoiceDate  UnitPrice  CustomerID      Country
415  2011-04-07 12:47:00      0.85     16701.0  United Kingdom
487  2011-10-20 18:24:00     49.95    13350.0  United Kingdom
613  2010-12-08 16:15:00     9.95     16252.0  United Kingdom
834  2011-07-19 12:26:00     4.15     12415.0  Australia
874  2011-02-01 15:38:00     1.45     17511.0  United Kingdom
...
162302 2011-07-01 13:06:00    12.75    17338.0  United Kingdom
162334 2011-04-20 13:39:00    2.55     17211.0  United Kingdom
162344 2011-05-12 19:43:00    4.95     13908.0  United Kingdom
162421 2011-02-01 15:38:00    2.55     17511.0  United Kingdom
162458 2011-02-07 14:46:00    8.50     18245.0  United Kingdom
```

[9288 rows x 8 columns]

```
# Deleting all the records which were returned
df=df[~df['InvoiceNo'].str.startswith('C',na=False)]
```

```
df.shape
```

(532621, 8)

```
# Checking for missing values
df.isnull().sum()
```

```
InvoiceNo          0
StockCode          0
Description       1454
Quantity          0
InvoiceDate       0
UnitPrice          0
CustomerID       134697
Country            0
dtype: int64
```

```
# Deleting all the missing values
df = df[pd.notnull(df['CustomerID'])]
```

```
# Creating a new column Total Sales
df['Total_sales'] = df['Quantity'] * df['UnitPrice']
```

```
df.describe().transpose()
```

	count	mean	std	min	25%
50% \					
Quantity	397924.0	13.021823	180.420210	1.0	2.00
6.00					
UnitPrice	397924.0	3.116174	22.096788	0.0	1.25
1.95					
CustomerID	397924.0	15294.315171	1713.169877	12346.0	13969.00

15159.00						
Total_sales	397924.0	22.394749	309.055588	0.0	4.68	
11.80						

Quantity	12.00	80995.00	75%	max
UnitPrice	3.75	8142.75		
CustomerID	16795.00	18287.00		
Total_sales	19.80	168469.60		

```
df[(df['Quantity'] > 500) & (df['Total_sales'] > 1000)].head(14)
```

	InvoiceNo	StockCode	Description
Quantity	\		
440	548011	22630	DOLLY GIRL LUNCH BOX
640			
4766	570094	22273	FELTCRAFT DOLL MOLLY
720			
6586	540815	21108	FAIRY CAKE FLANNEL ASSORTED COLOUR
3114			
9074	536809	84950	ASSORTED COLOUR T-LIGHT HOLDER
1824			
14761	567423	23285	PINK VINTAGE SPOT BEAKER
1944			
15029	554845	85123A	WHITE HANGING HEART T-LIGHT HOLDER
608			
15177	573008	84077	WORLD WAR 2 GLIDERS ASSTD DESIGNS
4800			
17773	574294	22086	PAPER CHAIN KIT 50'S CHRISTMAS
1020			
30612	558526	23170	REGENCY TEA PLATE ROSES
720			
36091	558554	23245	SET OF 3 REGENCY CAKE TINS
576			
37150	574294	21915	RED HARMONICA IN BOX
2100			
37246	543379	21623	VINTAGE UNION JACK MEMOBOARD
504			
38594	546789	85099B	JUMBO BAG RED RETROSPOT
1200			
43525	581483	23843	PAPER CRAFT , LITTLE BIRDIE
80995			

	InvoiceDate	UnitPrice	CustomerID	Country
Total_sales				
440	2011-03-29 11:14:00	1.65	14646.0	Netherlands
1056.00				
4766	2011-10-07 11:56:00	2.55	16029.0	United Kingdom
1836.00				
6586	2011-01-11 12:55:00	2.10	15749.0	United Kingdom

```

6539.40
9074 2010-12-02 16:48:00      0.55      15299.0 United Kingdom
1003.20
14761 2011-09-20 11:05:00      1.08      17450.0 United Kingdom
2099.52
15029 2011-05-26 19:49:00      3.24      17450.0 United Kingdom
1969.92
15177 2011-10-27 12:26:00      0.21      12901.0 United Kingdom
1008.00
17773 2011-11-03 15:47:00      2.55      16333.0 United Kingdom
2601.00
30612 2011-06-30 11:01:00      1.45      17949.0 United Kingdom
1044.00
36091 2011-06-30 12:50:00      4.15      17949.0 United Kingdom
2390.40
37150 2011-11-03 15:47:00      1.06      16333.0 United Kingdom
2226.00
37246 2011-02-07 15:37:00      6.38      18102.0 United Kingdom
3215.52
38594 2011-03-17 10:17:00      1.65      15769.0 United Kingdom
1980.00
43525 2011-12-09 09:15:00      2.08      16446.0 United Kingdom
168469.60

```

```
# Deleting record for Inovice no. 581483
df.drop(43525, axis =0, inplace = True)
```

```
df.describe().transpose()
```

	count	mean	std	min	25%
50% \Quantity	397922.0	12.818334	126.770315	1.0	2.00
6.00					
UnitPrice	397922.0	3.116180	22.096843	0.0	1.25
1.95					
CustomerID	397922.0	15294.310475	1713.172833	12346.0	13969.00
15159.00					
Total_sales	397922.0	21.971467	155.592056	0.0	4.68
11.80					

	75%	max
Quantity	12.00	74215.00
UnitPrice	3.75	8142.75
CustomerID	16795.00	18287.00
Total_sales	19.80	77183.60

```
df[df['Quantity'] == 74215]
```

	InvoiceNo	StockCode	Description	
Quantity \322892	541431	23166	MEDIUM CERAMIC TOP STORAGE JAR	74215

```

      InvoiceDate  UnitPrice  CustomerID          Country
Total_sales
322892 2011-01-18 10:01:00        1.04      12346.0  United Kingdom
77183.6

# Removing the record for invoice number 541431
df.drop(322892, axis =0, inplace = True)

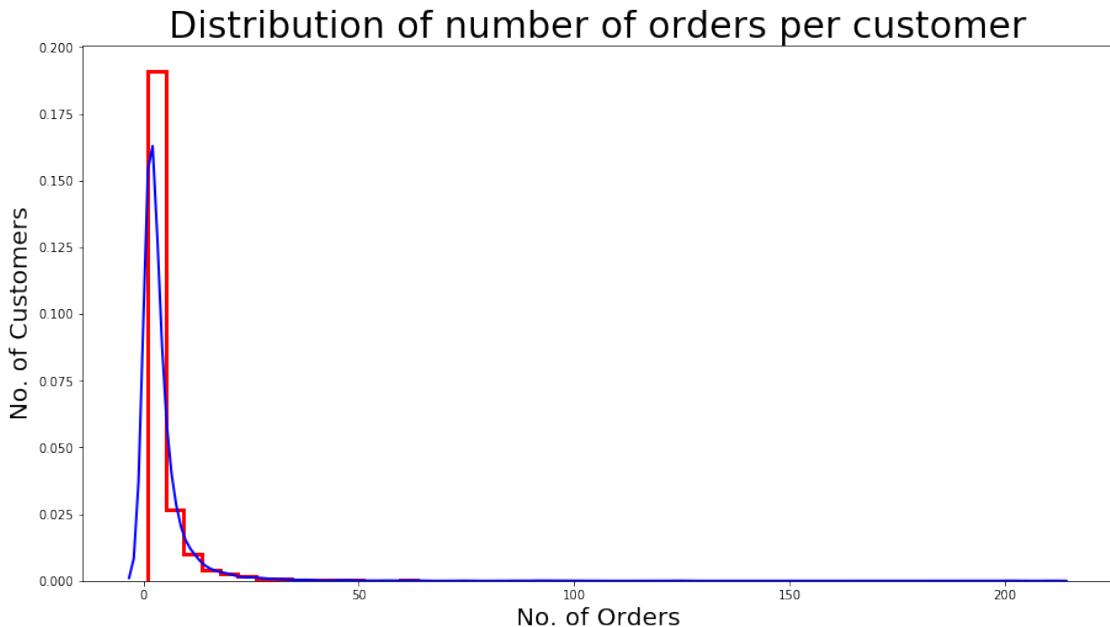
orders_n = df.groupby(['CustomerID']) ['InvoiceNo'].nunique()
more_than_one = np.sum(orders_n > 1) / df['CustomerID'].nunique()
print(f'{100 * more_than_one:.2f}% of customers ordered more than
once')

65.56% of customers ordered more than once

plt.figure(figsize=(15,8))
plt.style.use('_classic_test_patch')
px = sns.distplot(orders_n, hist = True, rug_kws={"color": "r"}, 
                  kde_kws={"color": "b", "lw": 2, "label": "KDE"}, 
                  hist_kws={"histtype": "step", "linewidth": 3, 
                            "alpha": 1, "color": "r"})

px.set_title('Distribution of number of orders per customer', fontsize
= 30)
px.set_xlabel('No. of Orders', fontsize = 20)
px.set_ylabel('No. of Customers', fontsize = 20);

```



```

order_country = df.groupby('Country').sum()
['Total_sales'].nlargest(10)
plt.figure(figsize=(30,12))

```

```

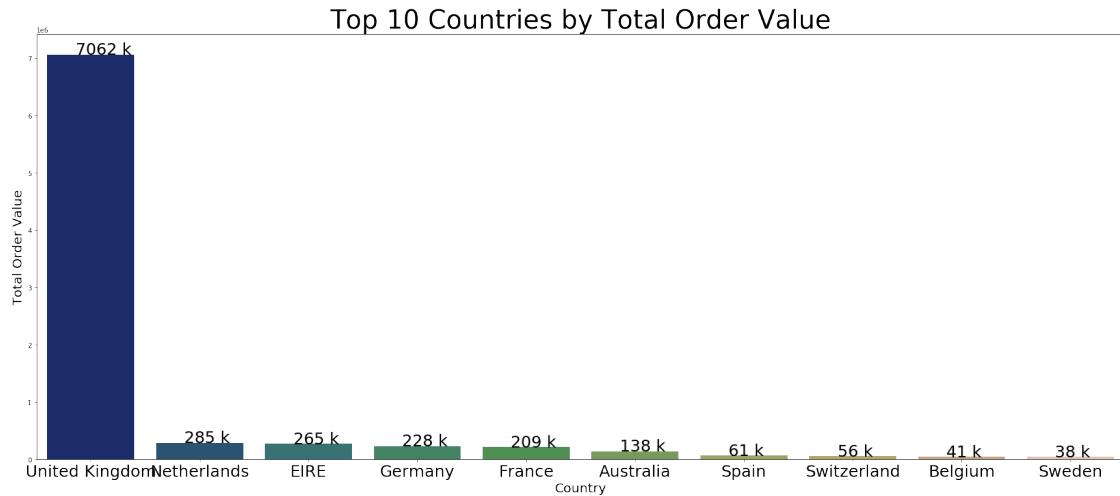
plt.style.use('_classic_test_patch')
px = sns.barplot(x = order_country.index, y = order_country.values,
palette = 'gist_earth')
px.set_xlabel('Country', fontsize = 20)
px.set_ylabel('Total Order Value', fontsize = 20)
px.set_title('Top 10 Countries by Total Order Value', fontsize = 40)

labels = [ '\n'.join(wrap(l, 15)) for l in order_country.index ]
px.set_xticklabels(labels, fontsize = 25)

value_ticks = []
for x in order_country.values:
    value_ticks.append(str(int(x/1000))+' k')
value_ticks[0]

for a, label in zip(px.patches, value_ticks):
    px.annotate(label, (a.get_x() + 0.26, a.get_height() + 2),
    fontsize = 25)

```



```

order_country = df.groupby('Country').sum()
['Total_sales'].nlargest(10)
plt.figure(figsize=(30,12))
plt.style.use('_classic_test_patch')
px = sns.barplot(x = order_country.index, y = order_country.values,
palette = 'gist_earth')
px.set_xlabel('Country', fontsize = 20)
px.set_ylabel('Total Order Value', fontsize = 20)
px.set_title('Bottom 10 Countries by Total Order Value', fontsize = 40)

labels = [ '\n'.join(wrap(l, 15)) for l in order_country.index ]
px.set_xticklabels(labels, fontsize = 25)

value_ticks = []
for x in order_country.values:

```

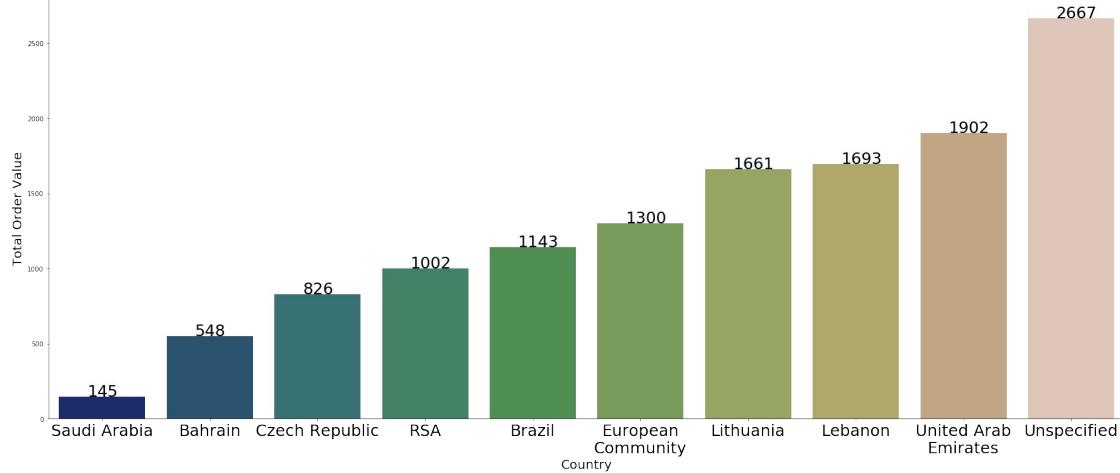
```

        value_ticks.append(str(int(x)))
value_ticks[0]

for a, label in zip(px.patches, value_ticks):
    px.annotate(label, (a.get_x() + 0.26, a.get_height() + 2),
    fontsize = 25)

```

Bottom 10 Countries by Total Order Value



```

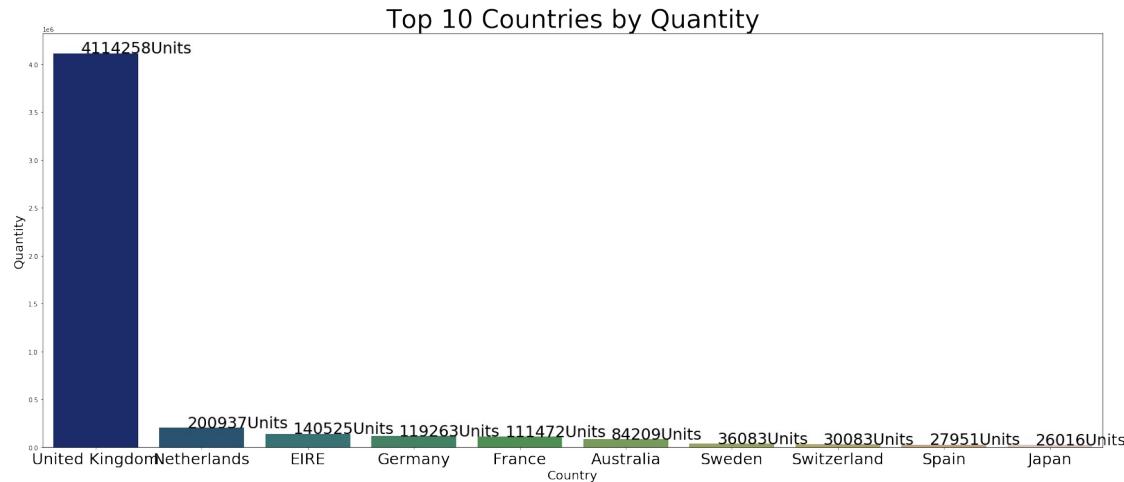
order_country =
df['Quantity'].groupby(df['Country']).agg('sum').sort_values(ascending
=False).head(10)
plt.figure(figsize=(30,12))
plt.style.use('_classic_test_patch')
px = sns.barplot(x = order_country.index, y = order_country.values,
palette = 'gist_earth')
px.set_xlabel('Country', fontsize = 20)
px.set_ylabel('Quantity', fontsize = 20)
px.set_title('Top 10 Countries by Quantity', fontsize = 40)

labels = [ '\n'.join(wrap(l, 15)) for l in order_country.index ]
px.set_xticklabels(labels, fontsize = 25)

value_ticks = []
for x in order_country.values:
    value_ticks.append(str(int(x))+'Units')
value_ticks[0]

for a, label in zip(px.patches, value_ticks):
    px.annotate(label, (a.get_x() + 0.26, a.get_height() + 2),
    fontsize = 25)

```



```

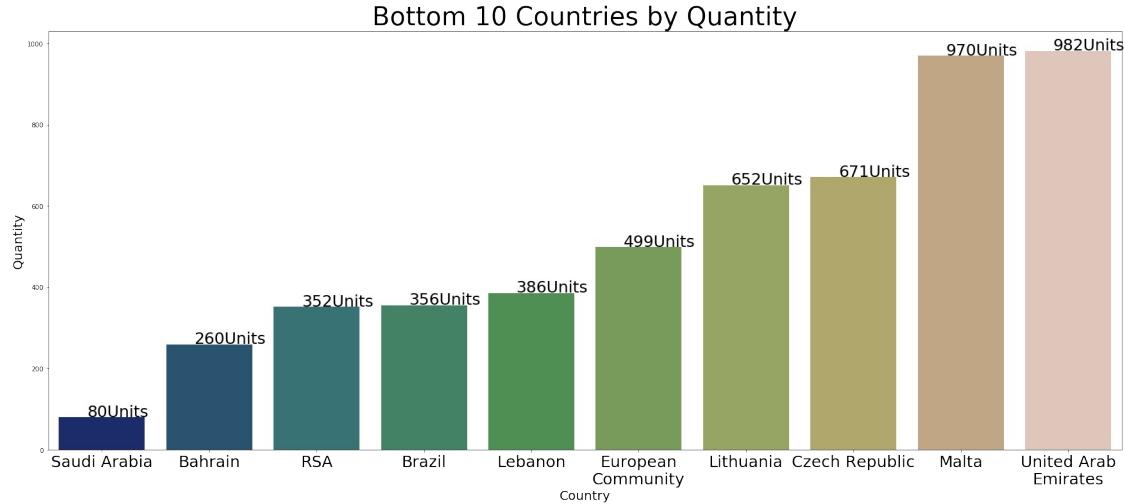
order_country =
df['Quantity'].groupby(df['Country']).agg('sum').sort_values(ascending
=True).head(10)
plt.figure(figsize=(30,12))
plt.style.use('_classic_test_patch')
px = sns.barplot(x = order_country.index, y = order_country.values,
palette = 'gist_earth')
px.set_xlabel('Country', fontsize = 20)
px.set_ylabel('Quantity', fontsize = 20)
px.set_title('Bottom 10 Countries by Quantity', fontsize = 40)

labels = [ '\n'.join(wrap(l, 15)) for l in order_country.index ]
px.set_xticklabels(labels, fontsize = 25)

value_ticks = []
for x in order_country.values:
    value_ticks.append(str(int(x))+'Units')
value_ticks[0]

for a, label in zip(px.patches, value_ticks):
    px.annotate(label, (a.get_x() + 0.26, a.get_height() + 2),
    fontsize = 25)

```



*#Total order value by country's*

```

sales_items = df.groupby('Description').sum()
['Total_sales'].nlargest(10)
plt.figure(figsize=(30,12))
plt.style.use('_classic_test_patch')
px = sns.barplot(x = sales_items.index, y = sales_items.values,
palette = 'gist_earth')
px.set_xlabel('Items', fontsize = 30)
px.set_ylabel('Total Order Value', fontsize = 30)
px.set_title('Top 10 Items by Total Sales', fontsize = 50)

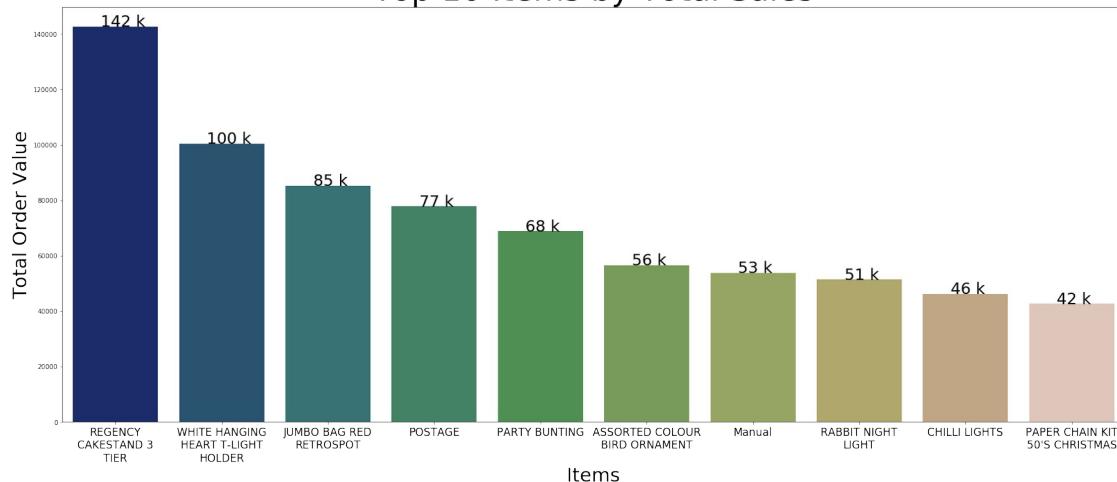
labels_x = [ '\n'.join(wrap(l, 15)) for l in sales_items.index ]
px.set_xticklabels(labels_x, fontsize = 17)

value_ticks = []
for x in sales_items.values:
    value_ticks.append(str(int(x/1000))+' k')

for a, label in zip(px.patches, value_ticks):
    px.annotate(label, (a.get_x() + 0.26, a.get_height() + 2),
    fontsize = 25)

```

## Top 10 Items by Total Sales



### #Cohort Analysis

```

# Creating a copy of the cleaned data
df_cohort = df.copy()

# As we have only 9 days information of December 2011, so we would not
# use that for cohort analysis.
df_cohort = df_cohort[~(df_cohort['InvoiceDate'] > '2011-12-01')]

# Extracting month and year column from the Invoice date column
df_cohort['invoice_month'] = df_cohort['InvoiceDate'].apply(lambda x:
x.strftime('%Y-%m'))

month_grp = df_cohort.groupby('invoice_month')

month_grp_agg = month_grp.agg({'CustomerID': pd.Series.nunique,
'Description': pd.Series.nunique,
'Quantity': np.sum, 'Total_sales':
np.sum})

month_grp_agg.rename(columns = {'CustomerID': '# customers',
'Description': 'tot_item_types',}, inplace = True)

```

month\_grp\_agg

invoice_month	# customers	tot_item_types	Quantity	Total_sales
2010-12	885	2418	312276	572705.490
2011-01	740	2117	274932	492261.440
2011-02	758	2122	265638	447137.350
2011-03	974	2240	348544	595500.760
2011-04	856	2225	292225	469200.361
2011-05	1056	2255	373685	678594.560
2011-06	991	2361	363699	661213.690
2011-07	949	2373	369432	600091.011

```

2011-08           935          2386   398938  645343.900
2011-09          1266          2581   544899  952838.382
2011-10          1364          2681   593908  1039318.790
2011-11          1665          2719   681888  1161817.380

```

*#Creating Label for customer cohort*

```

df_cohort.set_index('CustomerID', inplace = True)
df_cohort.head()

```

CustomerID	Quantity	InvoiceNo	StockCode	Description
16126.0	8	556072	20970	PINK FLORAL FELTCRAFT SHOULDER BAG
18118.0	2	551739	21559	STRAWBERRY LUNCH BOX WITH CUTLERY
15529.0	1	541658	21988	PACK OF 6 SKULL PAPER PLATES
14448.0	10	538364	85099C	JUMBO BAG BAROQUE BLACK WHITE
13911.0	4	552306	84789	ENCHANTED BIRD PLANT CAGE

CustomerID	InvoiceDate	UnitPrice	Country	Total_sales
16126.0	2011-06-08 14:57:00	3.75	United Kingdom	30.00
18118.0	2011-05-04 10:58:00	2.55	United Kingdom	5.10
15529.0	2011-01-20 12:16:00	0.85	United Kingdom	0.85
14448.0	2010-12-10 17:26:00	1.95	United Kingdom	19.50
13911.0	2011-05-08 15:20:00	3.75	United Kingdom	15.00

CustomerID	invoice_month
16126.0	2011-06
18118.0	2011-05
15529.0	2011-01
14448.0	2010-12
13911.0	2011-05

```

df_cohort['cust_cohort'] = df_cohort.groupby(level=0)[['InvoiceDate']].min().apply(lambda x: x.strftime('%Y-%m'))
df_cohort.head()

```

```

        InvoiceNo StockCode          Description
Quantity \\\
CustomerID

16126.0      556072      20970 PINK FLORAL FELTCRAFT SHOULDER BAG
8
18118.0      551739      21559 STRAWBERRY LUNCH BOX WITH CUTLERY
2
15529.0      541658      21988      PACK OF 6 SKULL PAPER PLATES
1
14448.0      538364      85099C JUMBO BAG BAROQUE BLACK WHITE
10
13911.0      552306      84789 ENCHANTED BIRD PLANT CAGE
4

        InvoiceDate UnitPrice          Country Total_sales
\ \\
CustomerID

16126.0      2011-06-08 14:57:00      3.75 United Kingdom      30.00
18118.0      2011-05-04 10:58:00      2.55 United Kingdom      5.10
15529.0      2011-01-20 12:16:00      0.85 United Kingdom      0.85
14448.0      2010-12-10 17:26:00      1.95 United Kingdom      19.50
13911.0      2011-05-08 15:20:00      3.75 United Kingdom      15.00

        invoice_month cust_cohort
CustomerID
16126.0      2011-06      2011-02
18118.0      2011-05      2010-12
15529.0      2011-01      2010-12
14448.0      2010-12      2010-12
13911.0      2011-05      2011-02

df_cohort.reset_index(inplace = True)
df_cohort.head()

        CustomerID InvoiceNo StockCode          Description
\ \\
0      16126.0      556072      20970 PINK FLORAL FELTCRAFT SHOULDER BAG
1      18118.0      551739      21559 STRAWBERRY LUNCH BOX WITH CUTLERY
2      15529.0      541658      21988      PACK OF 6 SKULL PAPER PLATES

```

```

3      14448.0      538364     85099C      JUMBO BAG BAROQUE BLACK WHITE
4      13911.0      552306     84789      ENCHANTED BIRD PLANT CAGE

```

	Quantity	InvoiceDate	UnitPrice	Country
Total_sales \ 30.00	8	2011-06-08 14:57:00	3.75	United Kingdom
5.10	2	2011-05-04 10:58:00	2.55	United Kingdom
0.85	1	2011-01-20 12:16:00	0.85	United Kingdom
19.50	10	2010-12-10 17:26:00	1.95	United Kingdom
4	4	2011-05-08 15:20:00	3.75	United Kingdom
15.00				

	invoice_month	cust_cohort
0	2011-06	2011-02
1	2011-05	2010-12
2	2011-01	2010-12
3	2010-12	2010-12
4	2011-05	2011-02

```
cust_grp = df_cohort.groupby(['cust_cohort','invoice_month'])
```

```
cust_grp_agg = cust_grp.agg({'CustomerID': pd.Series.nunique,
'Description': pd.Series.nunique,
'Quantity': np.sum, 'Total_sales': np.sum})
```

```
cust_grp_agg.rename(columns ={'CustomerID':'# customers',
'Description':'tot_item_types',}, inplace = True)
```

```
cust_grp_agg.head()
```

		# customers	tot_item_types	Quantity
Total_sales				
cust_cohort	invoice_month			
2010-12	2010-12	885	2418	312276
572705.49	2011-01	324	1787	158708
276237.69	2011-02	286	1697	136860
233845.37	2011-03	340	1834	177580
303119.39	2011-04	321	1781	131378
204407.66				

```

def cohort_num(df_cohort):
    df_cohort['cohort_num'] = np.arange(len(df_cohort)) + 1
    return df_cohort

cust_grp_agg = cust_grp_agg.groupby(level=0).apply(cohort_num)
cust_grp_agg.head(10)

# customers tot_item_types Quantity
Total_sales \
cust_cohort invoice_month

2010-12      2010-12          885       2418   312276
572705.49
276237.69      2011-01          324       1787   158708
233845.37      2011-02          286       1697   136860
303119.39      2011-03          340       1834   177580
204407.66      2011-04          321       1781   131378
336627.50      2011-05          352       1851   178599
314162.61      2011-06          321       1871   178645
310783.92      2011-07          309       1954   177188
331749.22      2011-08          313       1952   207217
472813.64      2011-09          350       2138   250114

            cohort_num
cust_cohort invoice_month
2010-12      2010-12          1
                    2011-01          2
                    2011-02          3
                    2011-03          4
                    2011-04          5
                    2011-05          6
                    2011-06          7
                    2011-07          8
                    2011-08          9
                    2011-09         10

```

#Project Task 2  
#Data Modelling

```

# Creating a copy of the original dataframe which will have records of
# December 2011 as well
rfm_df = df.copy(deep=True)

# Checking for duplicates
rfm_df.duplicated().sum()

5192

# Dropping all the duplicated records
rfm_df.drop_duplicates(inplace = True)

rfm_df.duplicated().sum()

0

rfm_df.shape

(392729, 9)

# The last transaction date we have is 12-09-2011, so we would take
# that as current date to calculate the recency
import datetime as dt
last_date = df['InvoiceDate'].max()

RFM_table = rfm_df.groupby('CustomerID').agg({'InvoiceDate': lambda x:
(last_date - x.max()).days, # Recency
                                                'InvoiceNo': 'count',
# Frequency
                                                'Total_sales': 'sum'})
# Monetary Value

RFM_table.rename(columns={'InvoiceDate': 'Recency',
                         'InvoiceNo': 'Frequency',
                         'Total_sales': 'Monetary'}, inplace=True)

RFM_table.head(10)

   Recency  Frequency  Monetary
CustomerID
12347.0        1       182    4310.00
12348.0       74       31    1797.24
12349.0       18       73    1757.55
12350.0      309       17    334.40
12352.0       35       85    2506.04
12353.0      203        4     89.00
12354.0      231       58   1079.40
12355.0      213       13    459.40
12356.0       22       59   2811.43
12357.0       32      131    6207.67

RFM_table.reset_index(inplace = True)

```

```

quartiles = RFM_table.quantile(q=[0.25,0.50,0.75])
quartiles = quartiles.to_dict()

RFM_table_copy = RFM_table.copy()

def RScore(x,p,d):
    if x <= d[p][0.25]:
        return 4
    elif x <= d[p][0.50]:
        return 3
    elif x <= d[p][0.75]:
        return 2
    else:
        return 1

def FMScore(x,p,d):
    if x <= d[p][0.25]:
        return 1
    elif x <= d[p][0.50]:
        return 2
    elif x <= d[p][0.75]:
        return 3
    else:
        return 4

RFM_table_copy['R'] = RFM_table_copy['Recency'].apply(RScore,
args=('Recency',quartiles))
RFM_table_copy['F'] = RFM_table_copy['Frequency'].apply(FMScore,
args=('Frequency',quartiles))
RFM_table_copy['M'] = RFM_table_copy['Monetary'].apply(FMScore,
args=('Monetary',quartiles))

RFM_table_copy['RFM'] = RFM_table_copy.R.map(str)
+RFM_table_copy.F.map(str)+RFM_table_copy.M.map(str)

RFM_table_copy['RFM_Score'] =
RFM_table_copy[['R','F','M']].sum(axis=1)
RFM_table_copy.head()

   CustomerID  Recency  Frequency  Monetary      R      F      M     RFM  RFM_Score
0      12347.0       1        182   4310.00    4      4      4    444         12
1      12348.0      74        31   1797.24    2      2      4    224          8
2      12349.0      18        73   1757.55    3      3      4    334         10
3      12350.0     309        17   334.40     1      1      2    112          4
4      12352.0      35        85   2506.04    3      3      4    334         10

RFM_table_copy.groupby('RFM_Score').agg({'Recency':
'mean','Frequency': 'mean', 'Monetary': ['mean', 'count'] }).round(1)

      Recency  Frequency  Monetary
           mean        mean      mean  count
RFM_Score

```

```

3      259.6      8.2     157.0    382
4      176.2     13.6     240.0    388
5      152.1     21.2     366.7    517
6      94.3      27.9     649.7    455
7      78.6      37.9     759.1    465
8      63.0      55.9     986.6    455
9      45.1      78.9    1391.5    413
10     31.5     110.6    2057.4    425
11     20.3     186.7    4055.8    388
12      6.2     367.9    9285.9    450

# Setting up the label for each client and adding the column "Label" to the dataframe

Segment = [0] * len(RFM_table_copy)

for i in range(0,len(RFM_table_copy)):

    if RFM_table_copy['RFM'][i] == '444':
        Segment[i] = "Platinum"

    elif RFM_table_copy['RFM'][i] == '413':
        Segment[i] = "Gold"

    elif RFM_table_copy['RFM'][i] == '323':
        Segment[i] = "Silver"

    else:
        Segment[i] = "Bronze"

RFM_table_copy['Segment'] = Segment

plt.figure(figsize=(20,10))
plt.style.use('_classic_test_patch')
ax = sns.countplot(x='Segment', data=RFM_table_copy,
palette="terrain",
order=RFM_table_copy['Segment'].value_counts().index[:5])
ax.set_ylabel('Count', fontsize = 25)
ax.set_xlabel('Customer Labels', fontsize = 35)
#ax.set_xticklabels(rfm_table['Label'].index, fontsize = 25)
ax.set_title('Customer Segmentation', fontsize = 38);

#labels_x = [ '\n'.join(wrap(l, 15)) for l in rfm_table['Label'].index ]
#ax.set_xticklabels(labels_x, fontsize = 19)

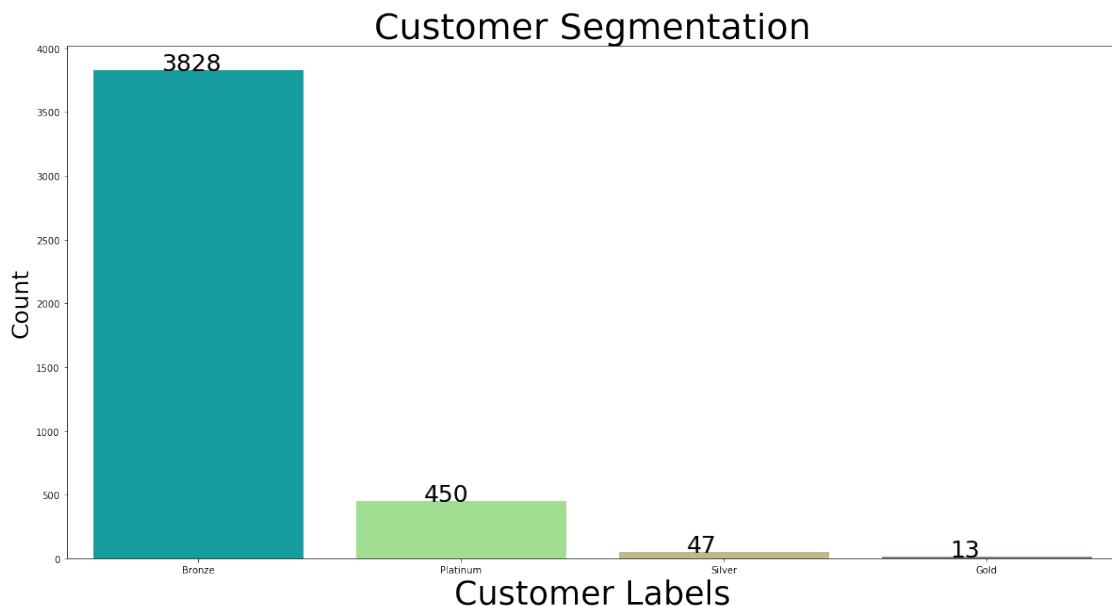
value_ticks = []
for x in RFM_table_copy['Segment'].value_counts():
    value_ticks.append(str(int(x)))

```

```

for a, label in zip(ax.patches, value_ticks):
    ax.annotate(label, (a.get_x() + 0.26, a.get_height() + 2),
    fontsize = 25)

```



```

RFM_table_copy.groupby('Segment').agg({'Recency':'mean',
'Frequency':'mean', 'Monetary':['mean','count']}).round(1)

```

Segment	Recency	Frequency	Monetary	
	mean	mean	mean	count
Bronze	102.5	58.9	1150.8	3828
Gold	9.4	11.9	1081.5	13
Platinum	6.2	367.9	9285.9	450
Silver	33.3	31.8	923.4	47

```

def skewness(df_skew, column):
    skew = stats.skew(df_skew[column])
    skewtest = stats.skewtest(df_skew[column])
    plt.title('Distribution of ' + column)
    sns.distplot(df_skew[column])
    print("{}'s: Skew: {}, : {}".format(column, skew, skewtest))
    return

```

```

# Plot all 3 graphs together for summary findings
plt.figure(figsize=(9, 9))

```

```

plt.subplot(3, 1, 1)
skewness(RFM_table_copy, 'Recency')

```

```

plt.subplot(3, 1, 2)
skewness(RFM_table_copy, 'Frequency')

```

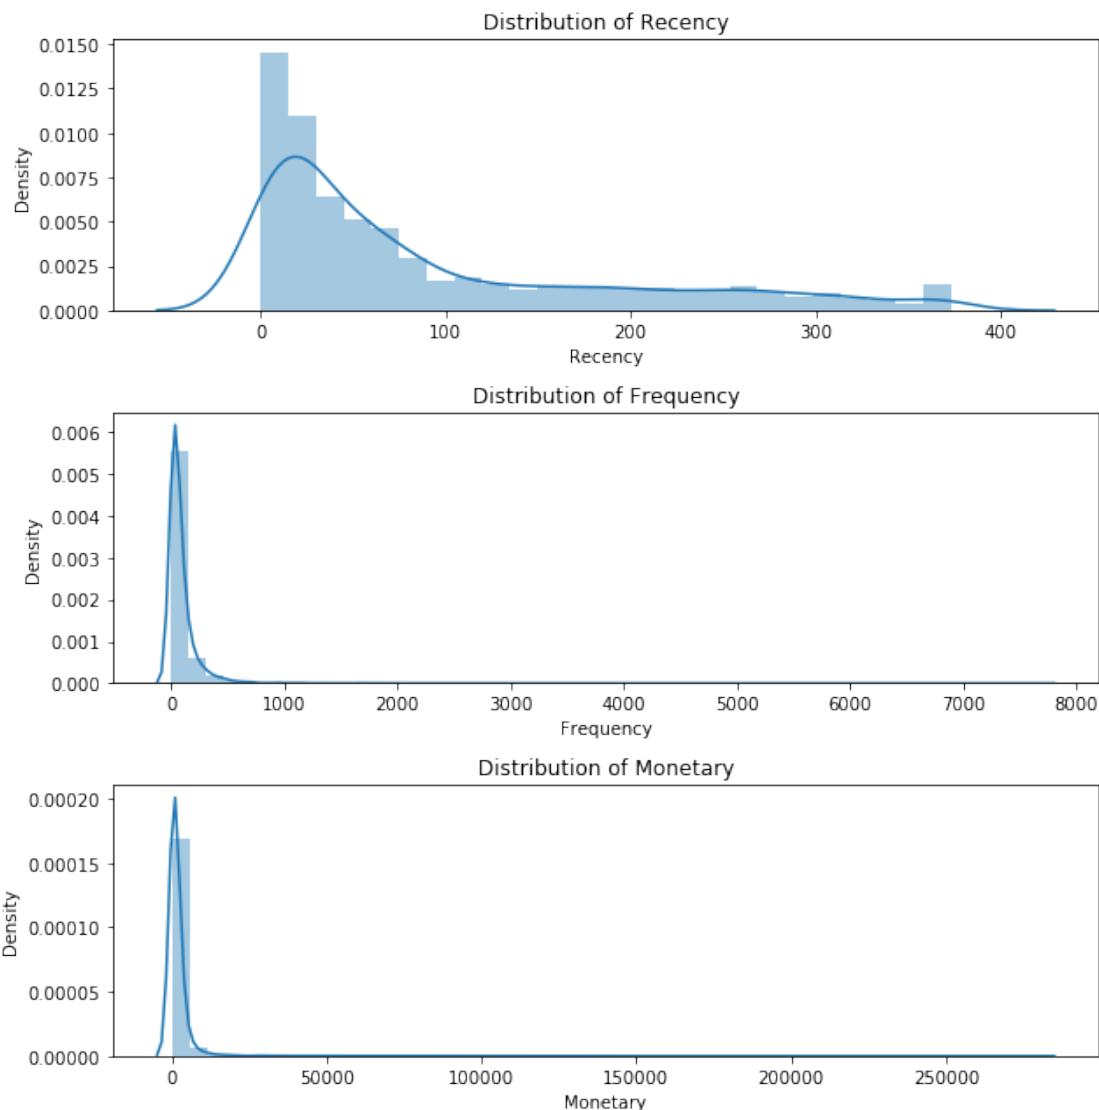
```

plt.subplot(3, 1, 3)
skewness(RFM_table_copy, 'Monetary')

plt.tight_layout()

Recency's: Skew: 1.2457125234954753, :
SkewtestResult(statistic=26.608212114006275,
pvalue=5.454166575151001e-156)
Frequency's: Skew: 18.036037104330713, :
SkewtestResult(statistic=83.48256629177025, pvalue=0.0)
Monetary's: Skew: 20.624815446376232, :
SkewtestResult(statistic=86.43722008800931, pvalue=0.0)

```



#Project Task 3  
#Data Modelling  
#K-means Clustering

```

df_rfm_log = RFM_table.copy()
print(df_rfm_log.describe())

   CustomerID      Recency      Frequency      Monetary
count    4338.000000  4338.000000  4338.000000  4338.000000
mean    15300.617796   91.511757   90.532273  1992.057929
std     1721.503870   99.963616  225.537197  8546.651225
min    12347.000000   0.000000   1.000000   0.000000
25%    13813.250000   17.000000  17.000000  306.210000
50%    15299.500000   50.000000  41.000000  668.040000
75%    16778.750000  141.000000  98.000000  1656.537500
max    18287.000000  373.000000  7676.000000 280206.020000

# A constant of 1 is added to the log to ensure we have all positive
values
df_rfm_log = np.log(df_rfm_log+1)
plt.figure(figsize=(9, 9))

plt.subplot(3, 1, 1)
skewness(df_rfm_log, 'Recency')

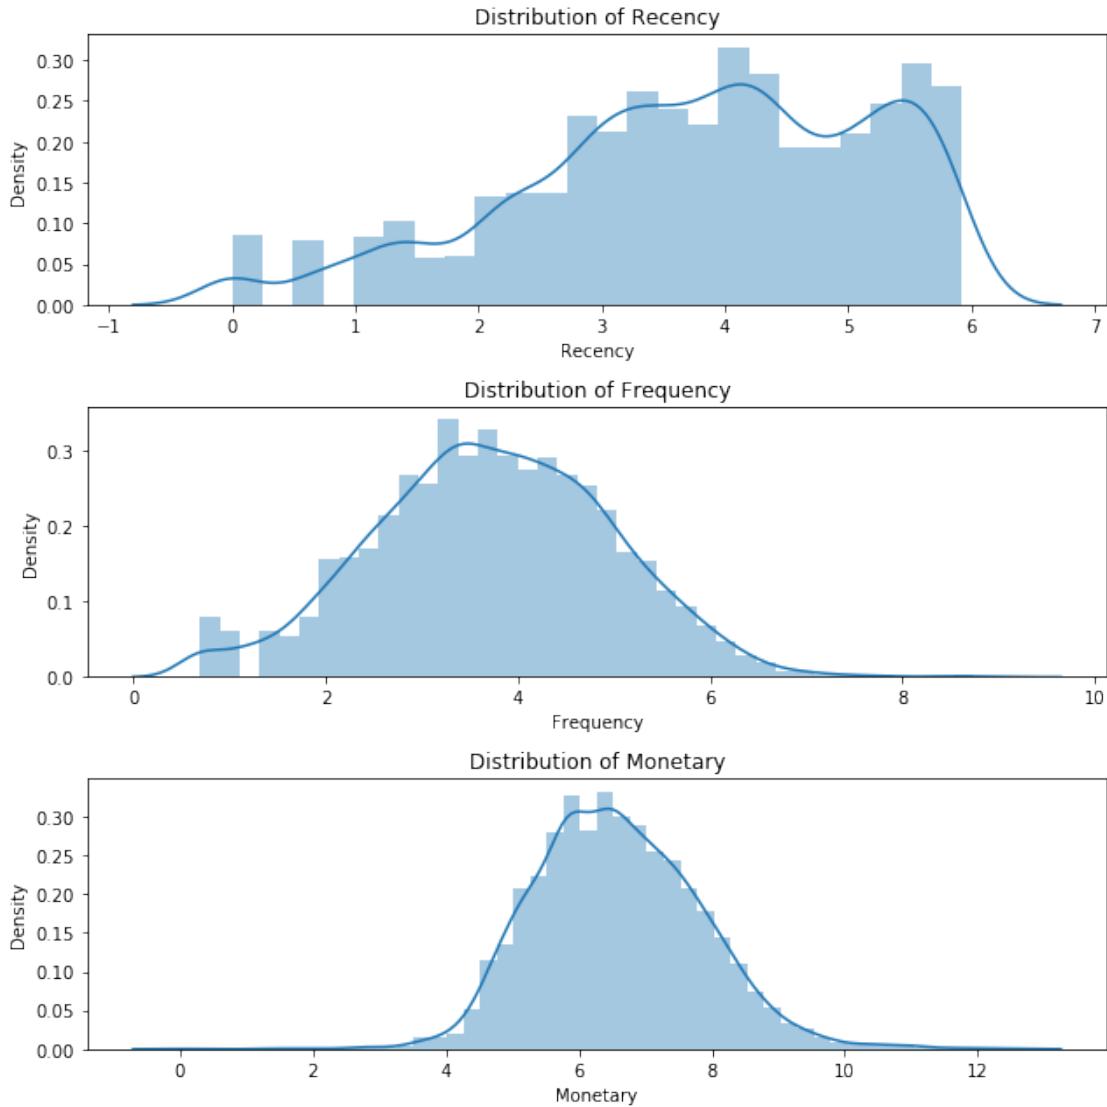
plt.subplot(3, 1, 2)
skewness(df_rfm_log, 'Frequency')

plt.subplot(3, 1, 3)
skewness(df_rfm_log, 'Monetary')

plt.tight_layout()

Recency's: Skew: -0.553047774132005, : SkewtestResult(statistic=-13.947408787843337, pvalue=3.262444980776586e-44)
Frequency's: Skew: -0.012391467138597751, : SkewtestResult(statistic=-0.3337508666486869, pvalue=0.7385675621614818)
Monetary's: Skew: 0.3271118898101863, :
SkewtestResult(statistic=8.59167125141112, pvalue=8.571333684337599e-18)

```



```
print(df_rfm_log.describe())
```

	CustomerID	Recency	Frequency	Monetary
count	4338.000000	4338.000000	4338.000000	4338.000000
mean	9.629314	3.768036	3.729132	6.583507
std	0.113513	1.431272	1.245280	1.260195
min	9.421249	0.000000	0.693147	0.000000
25%	9.533456	2.890372	2.890372	5.727531
50%	9.635641	3.931826	3.737670	6.505844
75%	9.727928	4.955827	4.595120	7.413088
max	9.814000	5.924256	8.945984	12.543284

```
# Standardizing the data
scaler = StandardScaler()
scaler.fit(df_rfm_log)
```

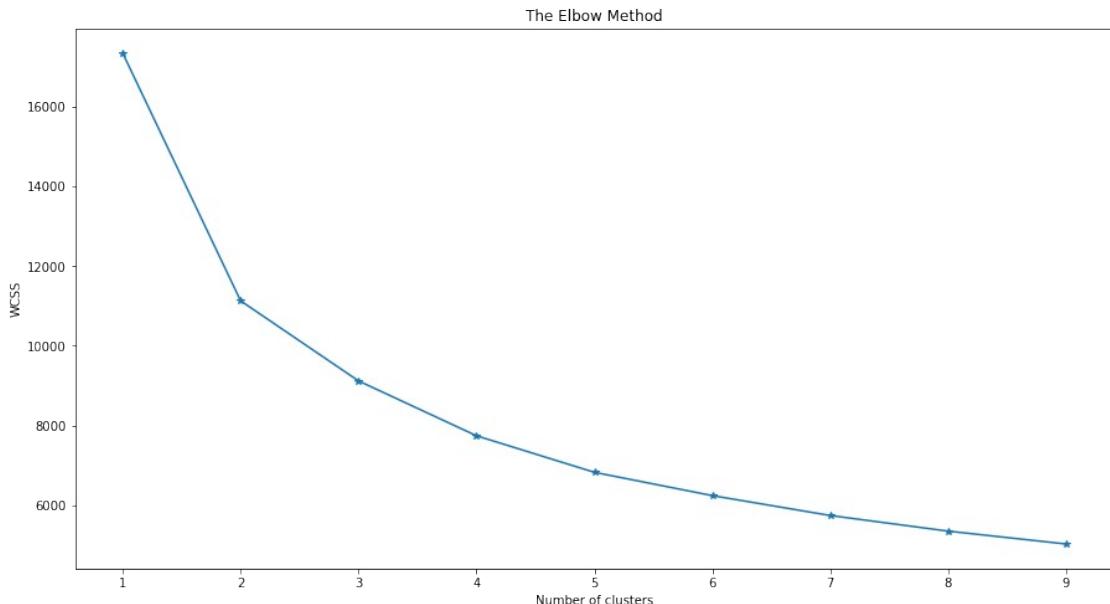
```

# Store it separately for clustering
rfm_scaled= scaler.transform(df_rfm_log)

wcss = []
for i in range(1, 10):
    kmeans = KMeans(n_clusters = i, init = 'k-means++', random_state =
42)
    kmeans.fit(rfm_scaled)
    wcss.append(kmeans.inertia_)

f, ax = plt.subplots(figsize=(15, 8))
plt.style.use('_classic_test_patch')
plt.plot(range(1, 10), wcss, '-*')
plt.title('The Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS');

```



```

def kmeans(rfm_scaled, clusters_number, original_df_rfm):

    kmeans = KMeans(n_clusters = clusters_number, random_state = 1)
    kmeans.fit(rfm_scaled)

    # Extract cluster labels
    cluster_labels = kmeans.labels_

    # Create a cluster label column in original dataset
    df_new = original_df_rfm.assign(Cluster = cluster_labels)

    # Initialise TSNE
    model = TSNE(random_state=1)
    transformed = model.fit_transform(df_new)

```

```
# Plot t-SNE
plt.title('Flattened Graph of {} Clusters'.format(clusters_number))
sns.scatterplot(x=transformed[:,0], y=transformed[:,1],
hue=cluster_labels, style=cluster_labels, palette="Set1")

return df_new

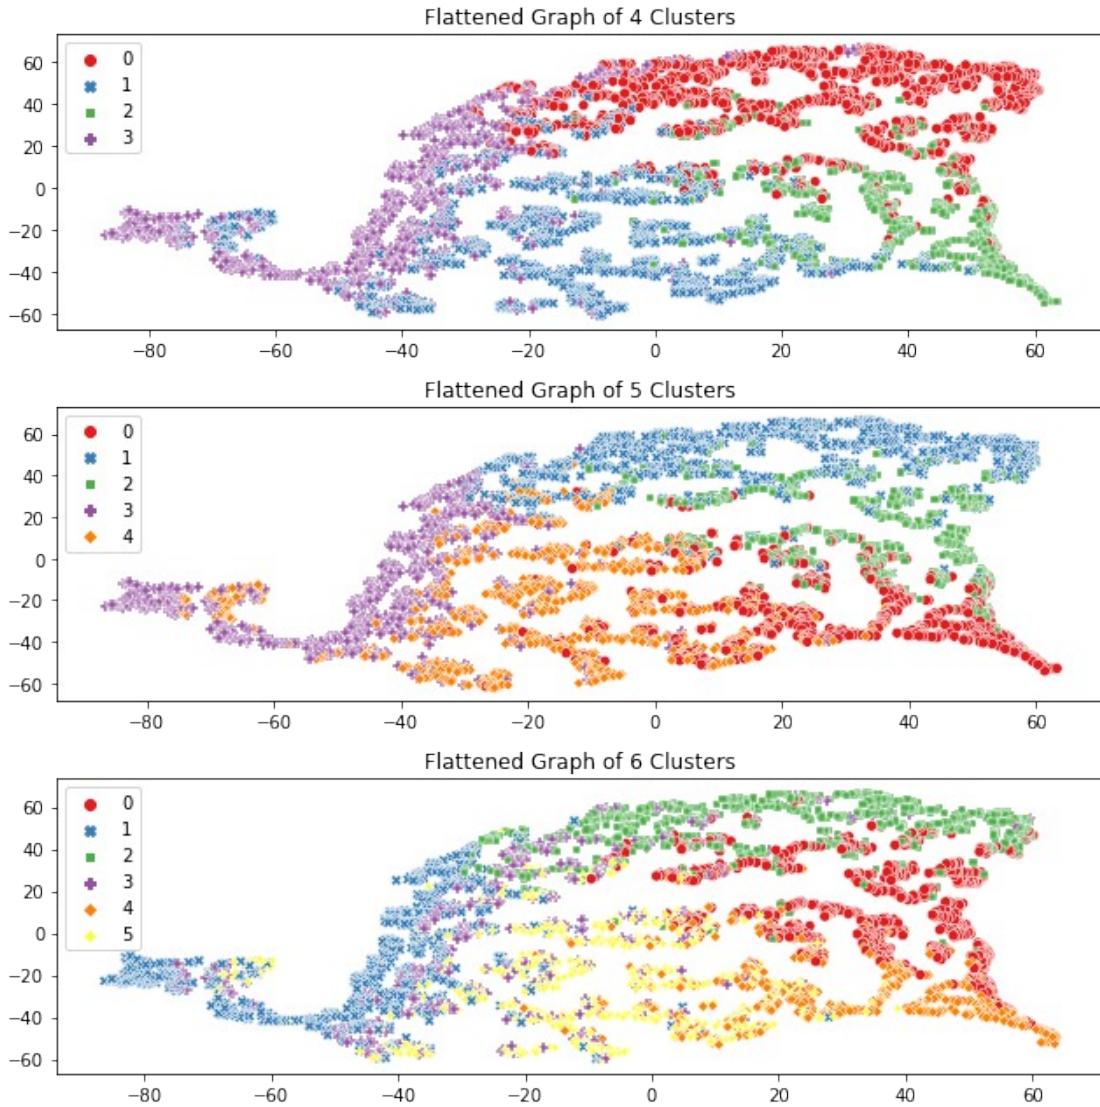
plt.figure(figsize=(9, 9))

plt.subplot(3, 1, 1)
df_rfm_k4 = kmeans(rfm_scaled, 4, RFM_table)

plt.subplot(3, 1, 2)
df_rfm_k5 = kmeans(rfm_scaled, 5, RFM_table)

plt.subplot(3, 1, 3)
df_rfm_k6 = kmeans(rfm_scaled, 6, RFM_table)

plt.tight_layout()
```



```

def wire_plot(rfm_scaled, df_rfm_kmeans, df_rfm_original):

    rfm_scaled = pd.DataFrame(rfm_scaled,
                               index=RFM_table.index,
                               columns=RFM_table.columns)
    rfm_scaled['Cluster'] = df_rfm_kmeans['Cluster']

    # Melt data into long format
    df_melt = pd.melt(rfm_scaled.reset_index(),
                      id_vars=['CustomerID', 'Cluster'],
                      value_vars=['Recency', 'Frequency',
'Monetary'],
                      var_name='Metric',
                      value_name='Value')

    plt.xlabel('Metric')

```

```
plt.ylabel('Value')
sns.pointplot(data=df_melt, x='Metric', y='Value', hue='Cluster')

return

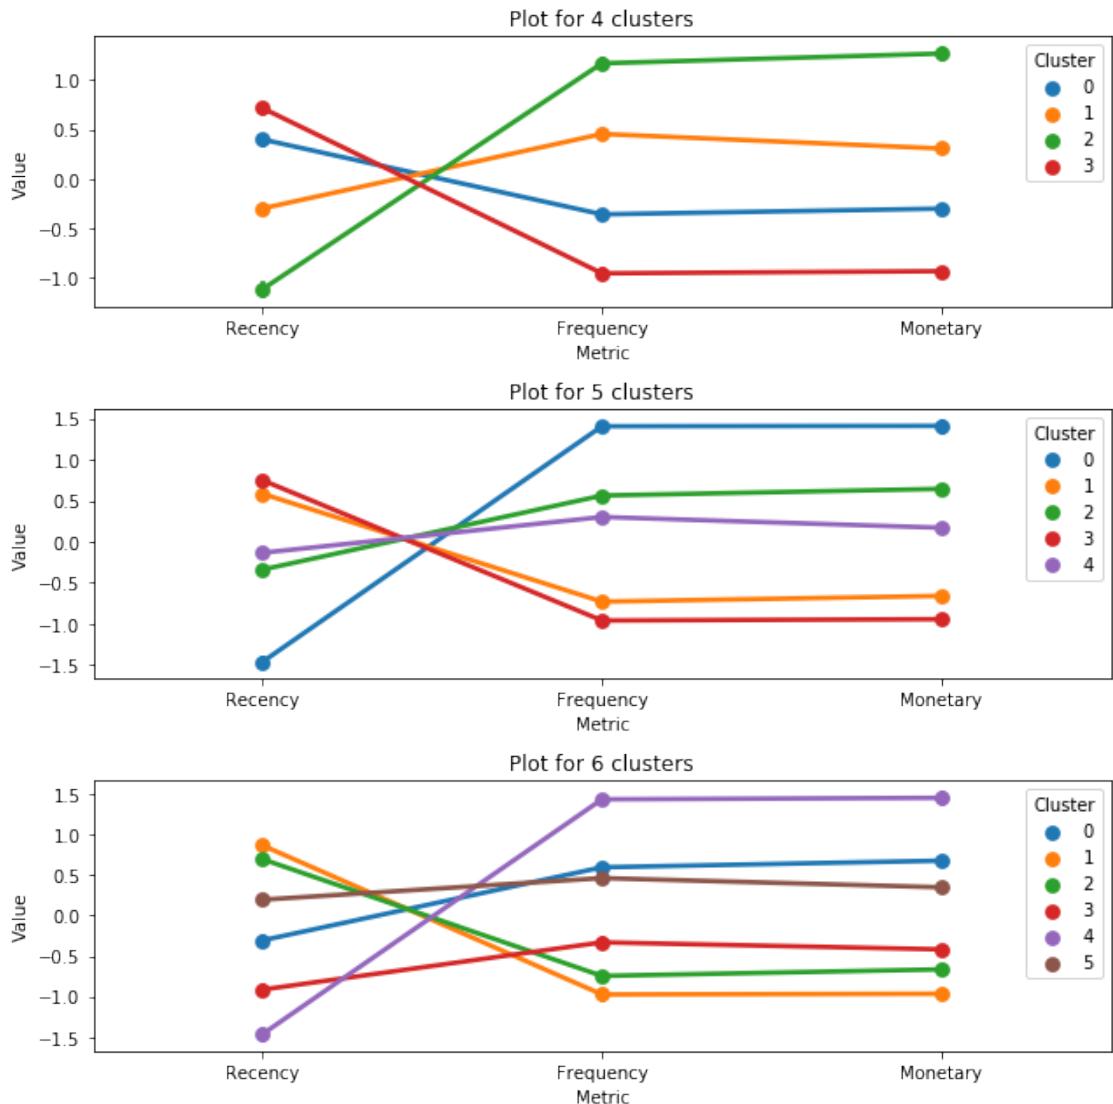
plt.figure(figsize=(9, 9))

plt.subplot(3, 1, 1)
plt.title('Plot for 4 clusters')
wire_plot(rfm_scaled, df_rfm_k4, RFM_table)

plt.subplot(3, 1, 2)
plt.title('Plot for 5 clusters')
wire_plot(rfm_scaled, df_rfm_k5, RFM_table)

plt.subplot(3, 1, 3)
plt.title('Plot for 6 clusters')
wire_plot(rfm_scaled, df_rfm_k6, RFM_table)

plt.tight_layout()
```



```
# clustering
kc = KMeans(n_clusters= 4, random_state=42)
kc.fit(rfm_scaled)

#Create a cluster label column in the original DataFrame
cluster_labels = kc.labels_

#Calculate average RFM values and size for each cluster:
df_rfm_k4 = df_rfm_log.assign(K_Cluster = cluster_labels)

#Calculate average RFM values and sizes for each cluster:
df_rfm_k4.groupby('K_Cluster').agg({'Recency': 'mean', 'Frequency' :
'mean',
'Monetary': [ 'mean',
'count'],}).round(0)
```

```

      Recency Frequency Monetary
      mean       mean      mean count
K_Cluster
0           4.0        3.0      6.0  1190
1           3.0        4.0      7.0  1195
2           5.0        3.0      5.0  1120
3           2.0        5.0      8.0   833

rfm_scaled =
pd.DataFrame(rfm_scaled,index=RFM_table.index,columns=RFM_table.columns)
rfm_scaled['K_Cluster'] = kc.labels_
rfm_scaled['Segment'] = RFM_table_copy['Segment']
rfm_scaled['CustomerID'] = RFM_table_copy['CustomerID']
rfm_scaled.reset_index(inplace = True)

#Melt the data into a long format so RFM values and metric names are stored in 1 column each
rfm_melt =
pd.melt(rfm_scaled,id_vars=['CustomerID','Segment','K_Cluster'],value_vars=['Recency', 'Frequency', 'Monetary'],
var_name='Metric',value_name='Value')
rfm_melt.head()

  CustomerID Segment K_Cluster Metric    Value
0     12347.0  Platinum         3 Recency -2.148608
1     12348.0    Bronze          0 Recency  0.383935
2     12349.0    Bronze          3 Recency -0.575496
3     12350.0    Bronze          0 Recency  1.375533
4     12352.0    Bronze          3 Recency -0.128933

f, (ax1, ax2) = plt.subplots(1,2, figsize=(15, 8))
sns.lineplot(x = 'Metric', y = 'Value', hue = 'Segment', data = rfm_melt,ax=ax1)

# a snake plot with K-Means
sns.lineplot(x = 'Metric', y = 'Value', hue = 'K_Cluster', data = rfm_melt,ax=ax2)

plt.suptitle("Plot of RFM Segments & Clusters", fontsize=24)
plt.show()

```

## Plot of RFM Segments & Clusters

